


Spam:

step size : .5

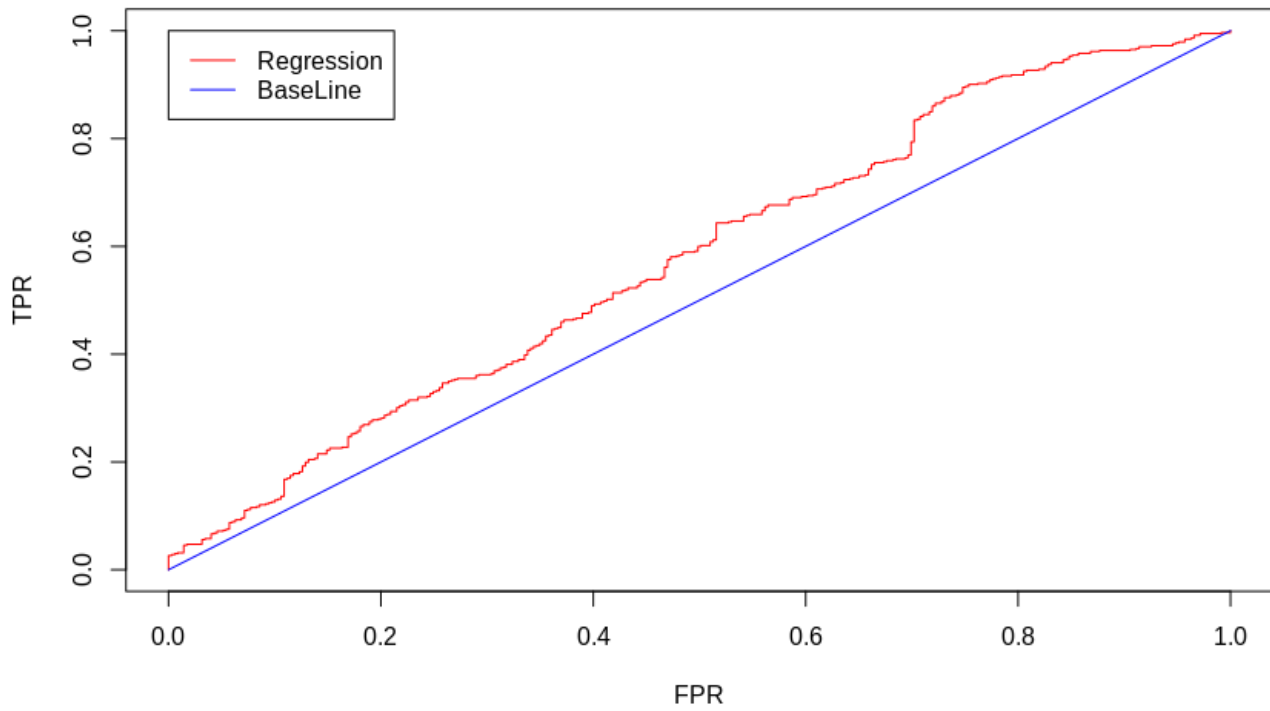
max iterations = 1500

Spam data, Training set Vs. Validation set, loss function over number of observations, N = 15



Obviously something went wrong, they almost seem like these are upside down. The train data is supposed to always go down and the validation set is supposed to go down and eventually go back up. However this graph seems to indicate that the appropriate number of iterations is 1500, which is not correct for a very simple algorithm like this.

ROC curve of linear regression Vs Baseline



This is the ROC curve for our function. As you can see it is not much better than the baseline.